Deep Learning Project

KYRIAKI BEI

The Project

Subject

Speaker Identification by Voice.

Dataset

VoxCeleb1: An audio dataset consisting of over 100,000 short clips of human speech for 1.251 celebrities, extracted from interview videos uploaded to YouTube.

	dev files
# speakers	1,211
# audio files	148,642

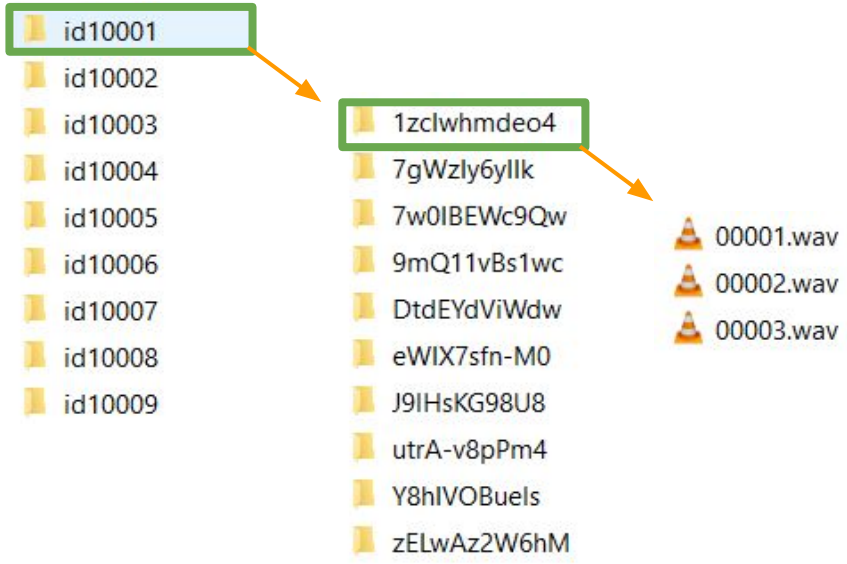

File Handling

From the **1,211** total speakers, **10** were selected for the classification task according to their number of audio files.

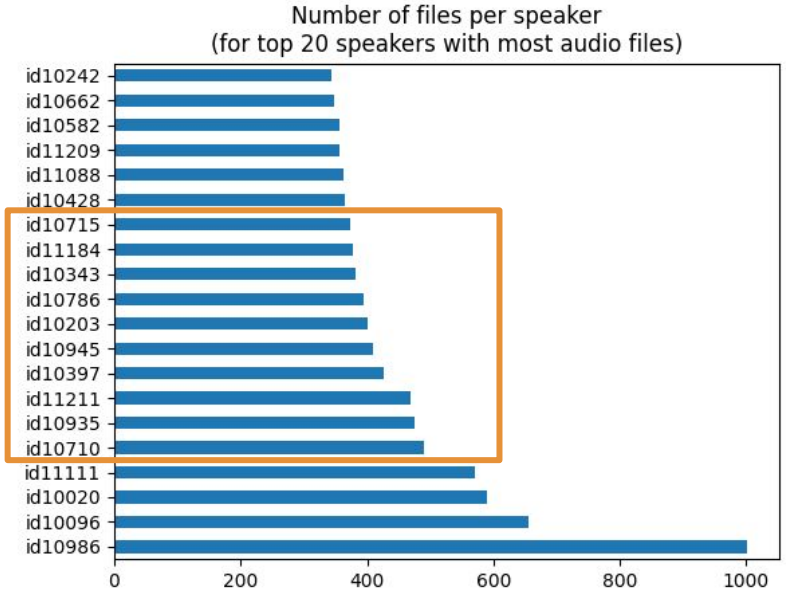

Dictionary that maps the 10 speakers' ids to their audio files


```
{
    "id10710": [
        "data\\voxceleb_data\\wav\\id10710\\230QPHWW7fM\\00001.wav",
        "data\\voxceleb_data\\wav\\id10710\\230QPHWW7fM\\00002.wav",
        "data\\voxceleb_data\\wav\\id10710\\230QPHWW7fM\\00003.wav",
        "data\\voxceleb_data\\wav\\id10710\\230QPHWW7fM\\00004.wav",
        "data\\voxceleb_data\\wav\\id10710\\230QPHWW7fM\\00005.wav",
        "data\\voxceleb_data\\wav\\id10935\\Z1x9xyh8NjA\\00009.wav"
    ],
    "id11211": [
        "data\\voxceleb_data\\wav\\id11211\\0LUUizW5U48\\00001.wav",
        "data\\voxceleb_data\\wav\\id11211\\0LUUizW5U48\\00002.wav",
        "data\\voxceleb_data\\wav\\id11211\\6mnL46vX6NY\\00001.wav",
        "data\\voxceleb_data\\wav\\id11211\\6mnL46vX6NY\\00002.wav",
```

The Dataset

The dataset used for classification consists of **10 speakers** and **4,195** audio files in total.

Selected speakers

1. David Attenborough
2. Gloria Steinem
3. J.J. Abrams
4. Louis C.K.
5. Lucie Arnaz
6. Meat Loaf
7. Peter Jackson
8. Quentin Tarantino
9. Tom Hooper
10. Vince Gilligan

Genders Distribution

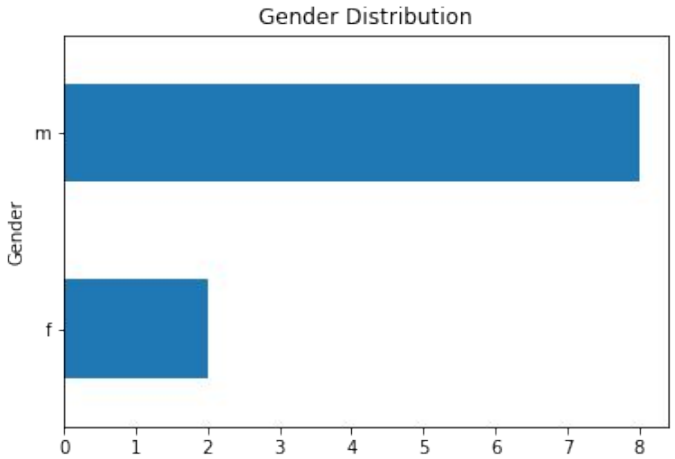

Nationalities Distribution

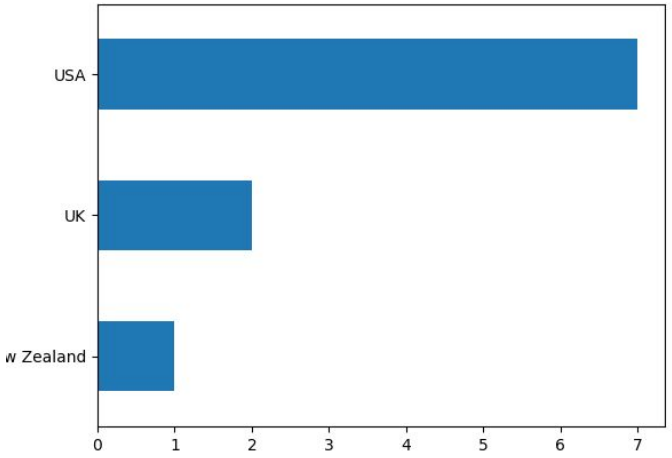

Workflow

STFT

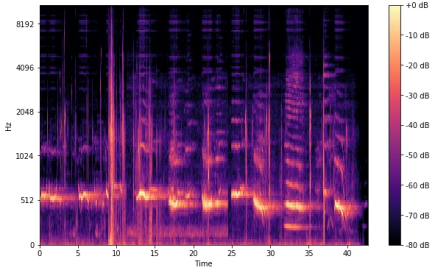

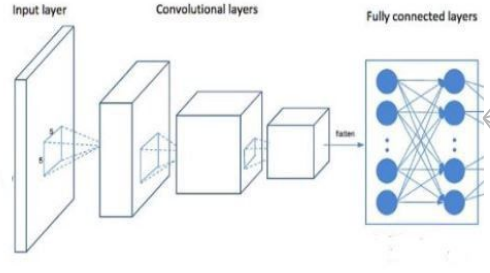

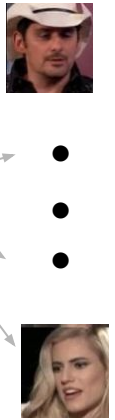

waveform

- same sample rate
- same # channels

mel spectrogram

CNN

FC layer

Audio Preprocessing

sample rate = 16000 kHz

mono signal

mel spectrogram

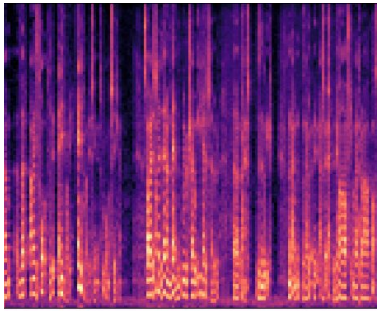

Audio Preprocessing

- optimal image width for spectrogram images found

  optimal width = 95th-percentile of image widths

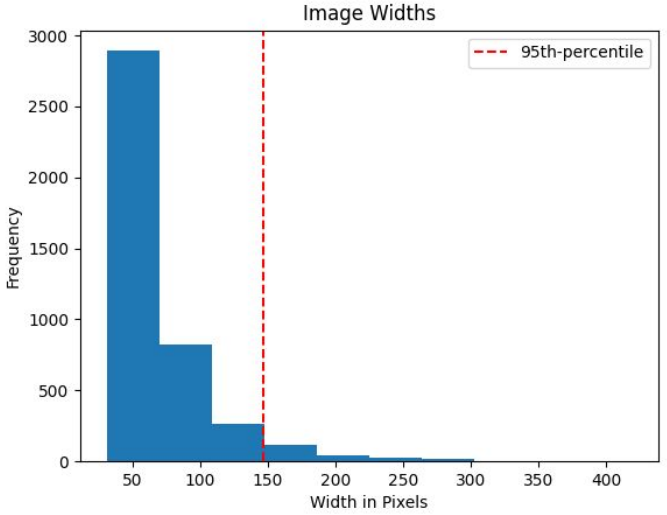

- crop or pad images to have width equal to the optimal image width

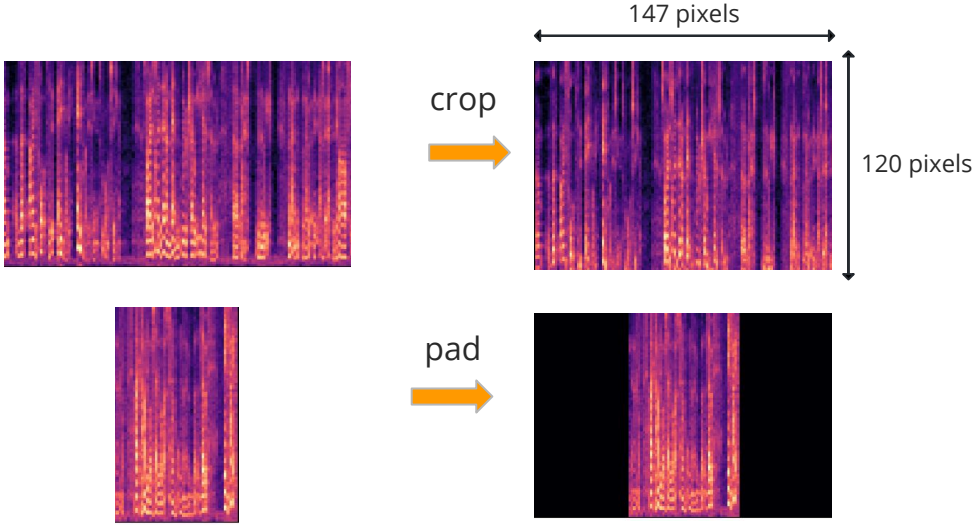

The CNN Model

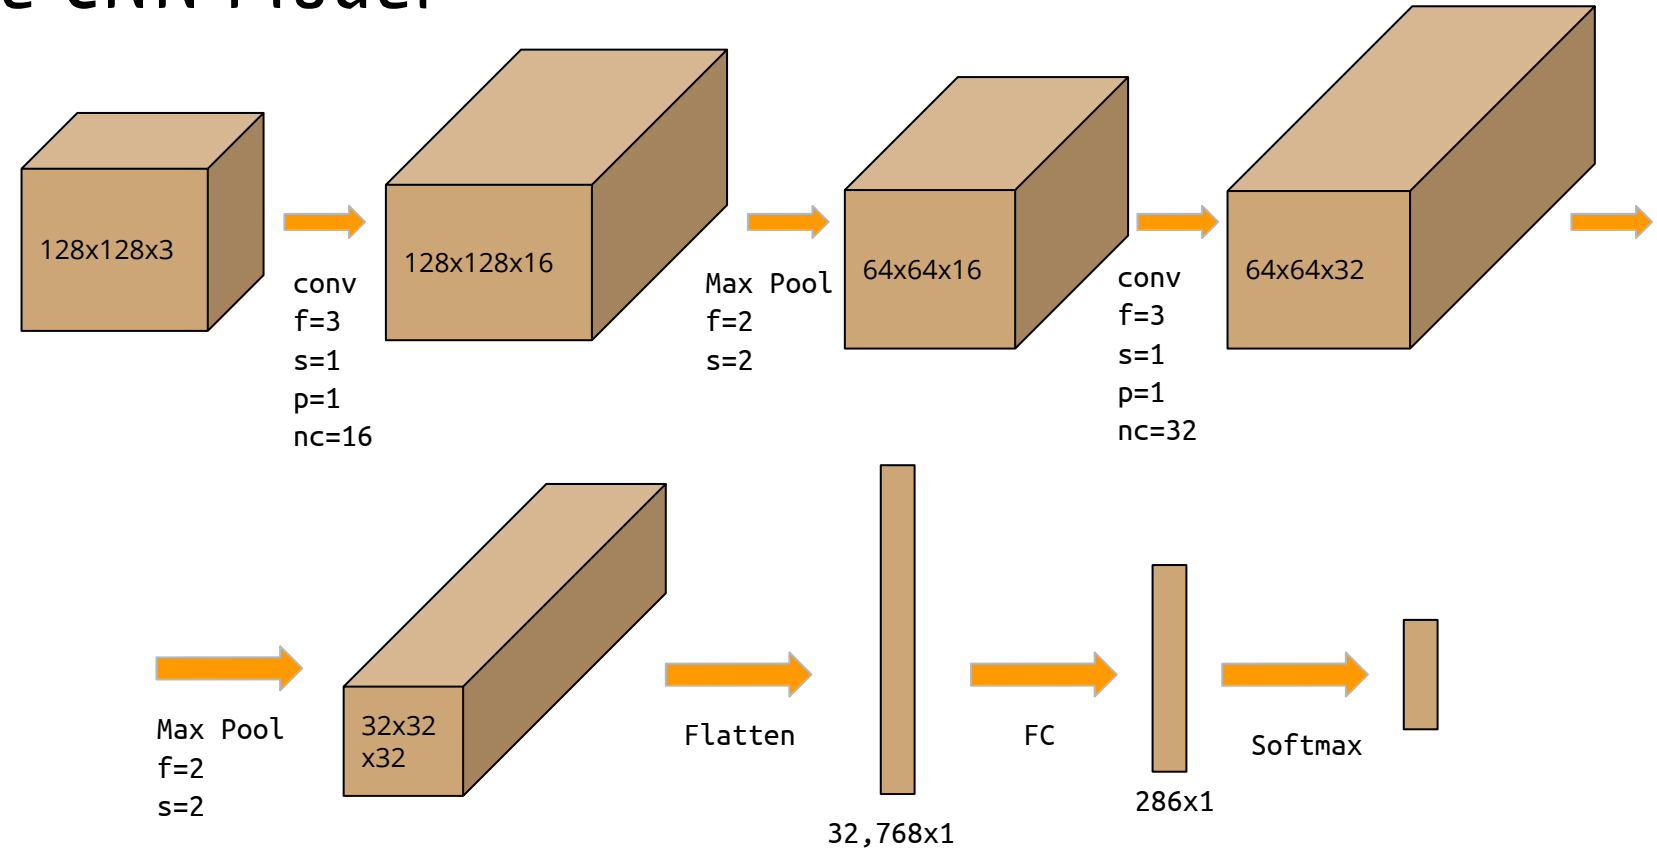

Results

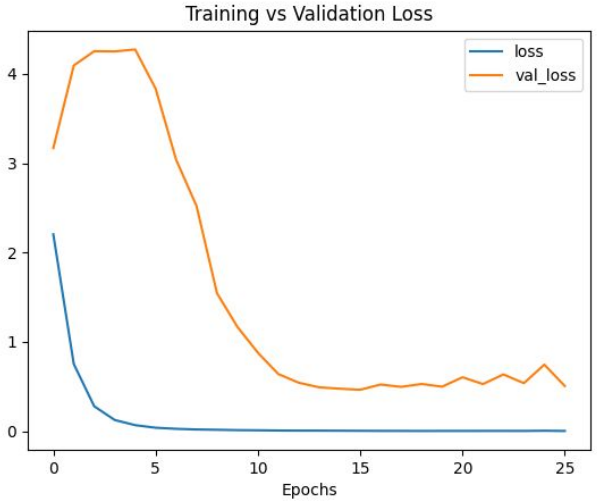

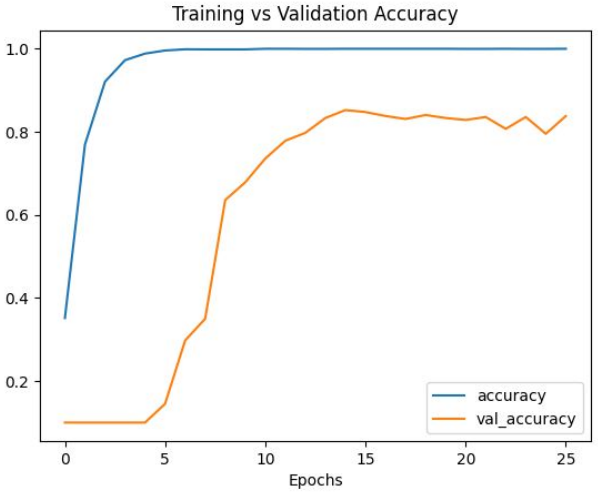

Test loss: 0.324
Test accuracy: 90.5 %

Some Predictions

- Predicted speaker **Gloria Steinem** with probability **99.8 %** (True speaker: Gloria Steinem)
- Predicted speaker **Tom Hooper** with probability **95.4 %** (True speaker: Tom Hooper)
- Predicted speaker **Quentin Tarantino** with probability **91.6 %** (True speaker: Quentin Tarantino)